

Community Use of XALT in Its First Year in Production

HUST 2015
Austin, TX



Reuben D. Budiardja

National Institute for Computational Sciences
The University of Tennessee

with Mark Fahey (ANL), Robert McLay (TACC),
Prasad Maddumage Don (FSU),
Bilel Hadri (KAUST), Doug James (TACC)



National Science Foundation
WHERE DISCOVERIES BEGIN

<https://github.com/Fahey-McLay/xalt>

Talk Outline

- Introduction to XALT

- Motivation
- How It Works

Getting Data Out of XALT

- Compilers, Libraries, Executables Usage Reports
- Other Use Cases

- New Functionality

- Function Tracking
- GUI (Web)-Based Reports
- User Software Provenance

Introduction to XALT

Motivation

Most computing center needs to answer the questions:

- How many users and projects use a particular library or executable ?

How many users use which compilers ?

- Which center provided packages are used often ? and which one are never used ?
- Which users or applications still use old version of certain library, compiler, or executable ?

Are there any widely used user-installed package that a center should provide instead ?

XALT is a tool to collect accurate, detailed, and continuous job-level and link-time data, and store them in a database.

XALT is a tool to collect accurate, detailed, and continuous job-level and link-time data, and store them in a database.

XALT collects information to answer questions on software usage


Goals

- ***Automatic, continuous*** census of libraries and applications
 - Collect job-level and link-time level data for subsequent analytics
- Must be transparent to user, avoid impacting the user experience
- Must work seamlessly on any system: workstation, cluster, high-end supercomputer
 - Must be a lightweight solution

Approach: Link-time Level

Intercept linker at link-time:



- Wrap the (GNU) linker (ld) and parse the command line
 - Capture only the object files actually linked with the executable
- Stores the results using a chosen *transmission style*
- Insert an XALT's ELF section header to the executable



Approach: Execution-time Level

Intercept job launcher to get execution environment:



- Wrap job-launcher (aprun, ibrun, mpirun, ...) with a corresponding script
- Extract previously inserted XALT's ELF header (if any)
- Extract environment variables
 - Job-specific environment (e.g. PBS_JOBID, etc)
 - Dynamics libraries loaded at runtime
 - Record job start and end time

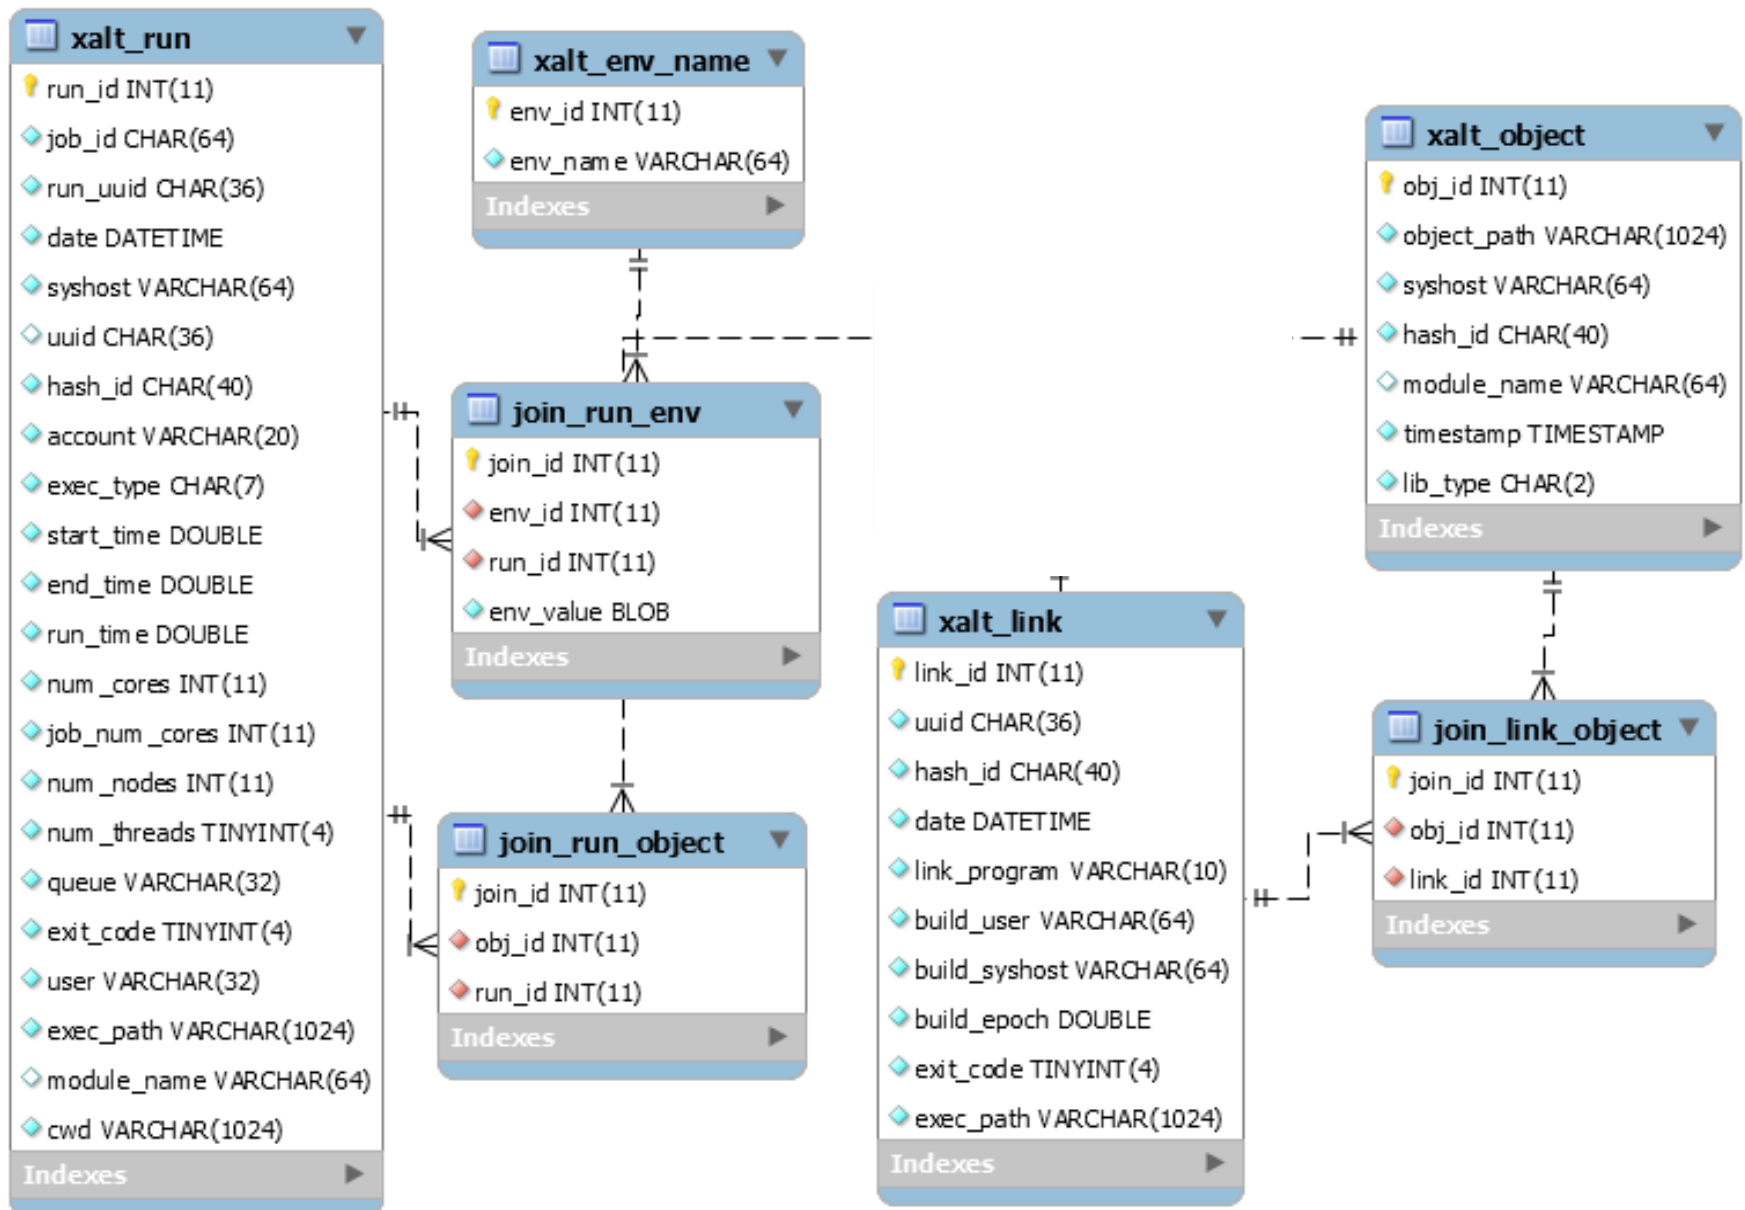


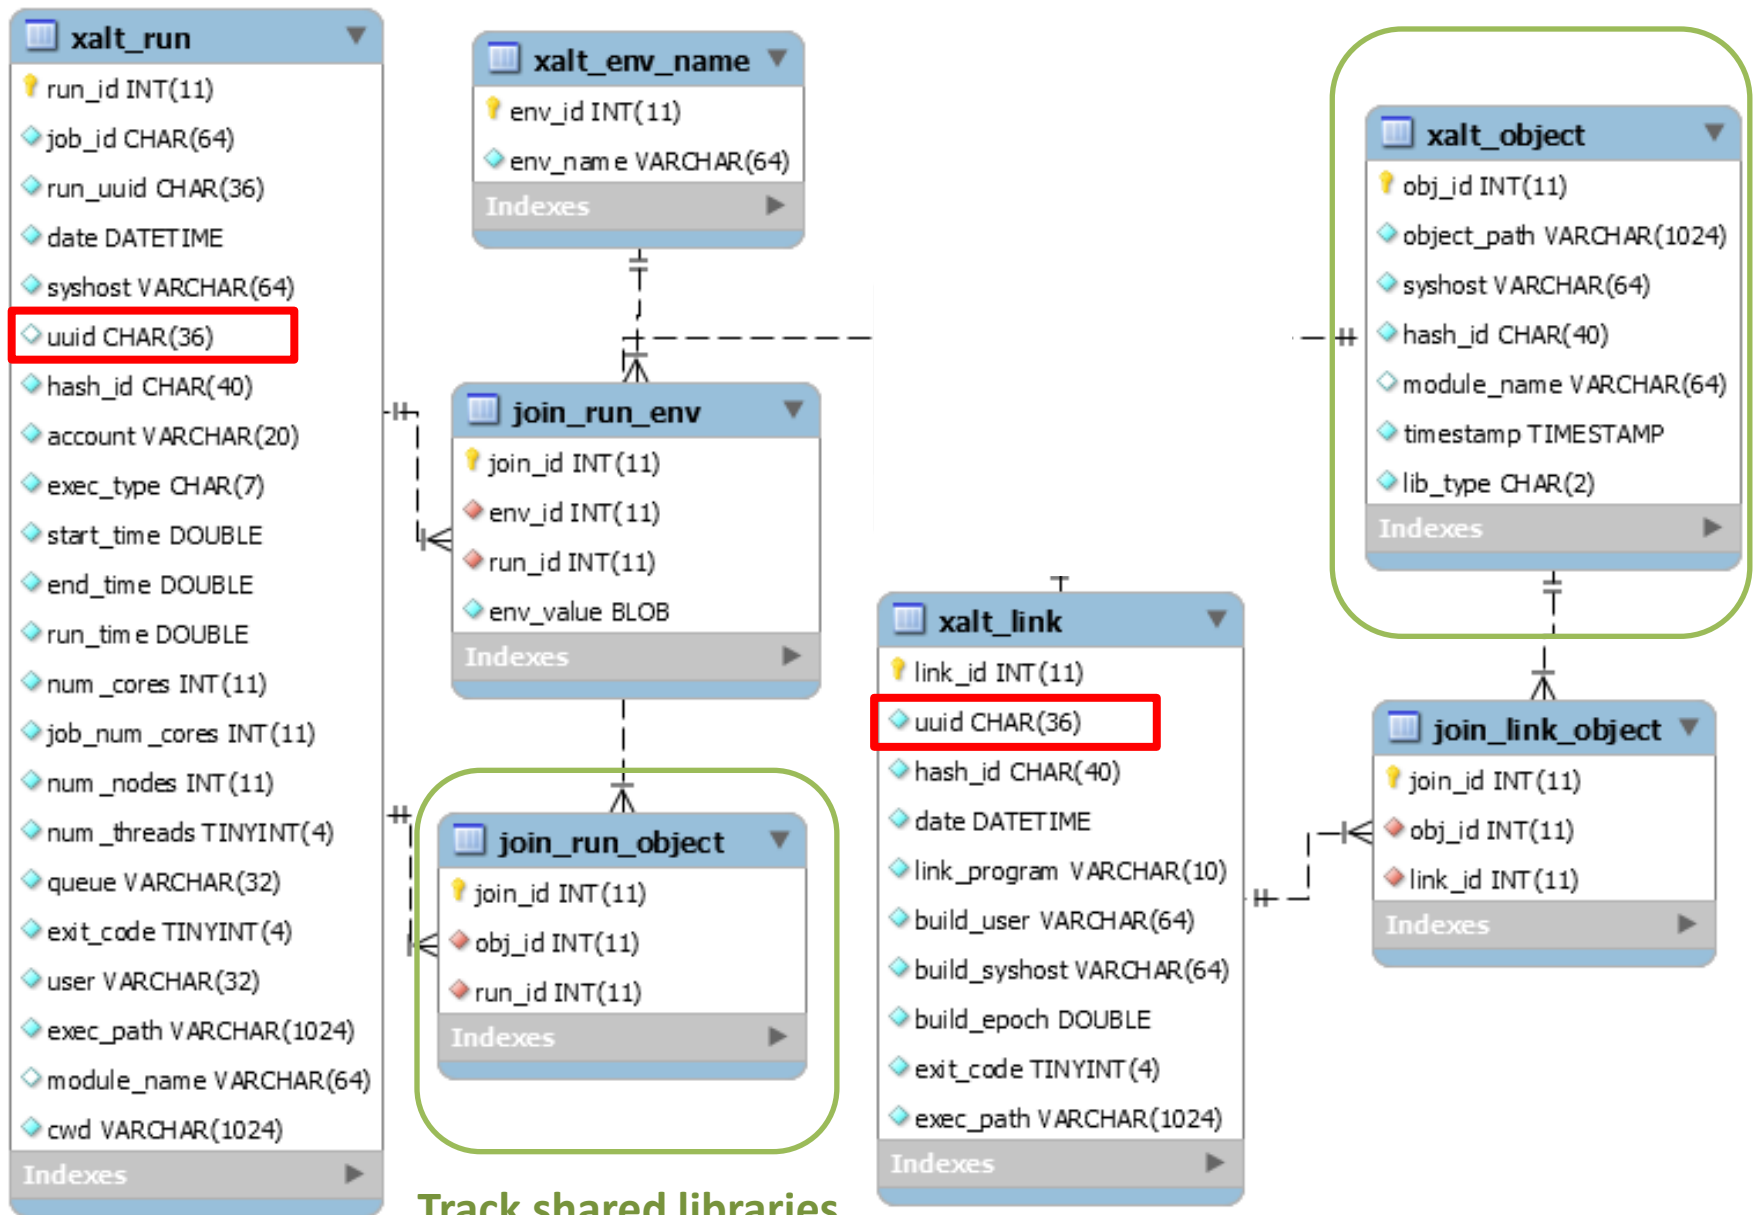
xalt_link	
link_id	INT(11)
uuid	CHAR(36)
hash_id	CHAR(40)
date	DATETIME
link_program	VARCHAR(10)
build_user	VARCHAR(64)
build_syshost	VARCHAR(64)
build_epoch	DOUBLE
exit_code	TINYINT(4)
exec_path	VARCHAR(1024)
Indexes	

xalt_object	
obj_id	INT(11)
object_path	VARCHAR(1024)
syshost	VARCHAR(64)
hash_id	CHAR(40)
module_name	VARCHAR(64)
timestamp	TIMESTAMP
lib_type	CHAR(2)
Indexes	

join_link_object	
join_id	INT(11)
obj_id	INT(11)
link_id	INT(11)
Indexes	









Getting Data Out of XALT

Community Usage Reports

Compiler Usage


- XALT stores “link program”: the program that calls the linker
 - A proxy for the compiler → main() compiler
 - Will miss mixed language compilation
- Can associate “compiler” with every linking event

Compiler Usage on Darter



```
SELECT link_program , count(*) as count
FROM xalt_link
WHERE build_syshost = [syshost]
GROUP BY link_program ORDER BY count desc
```


Compiler Usage Ratio per User




Is there a way to tell if someone used a compiler once (or a little), before giving up ?

Compiler Usage: TACC, FSU, KAUST


TACC



FSU



KAUST




Most Used Libraries

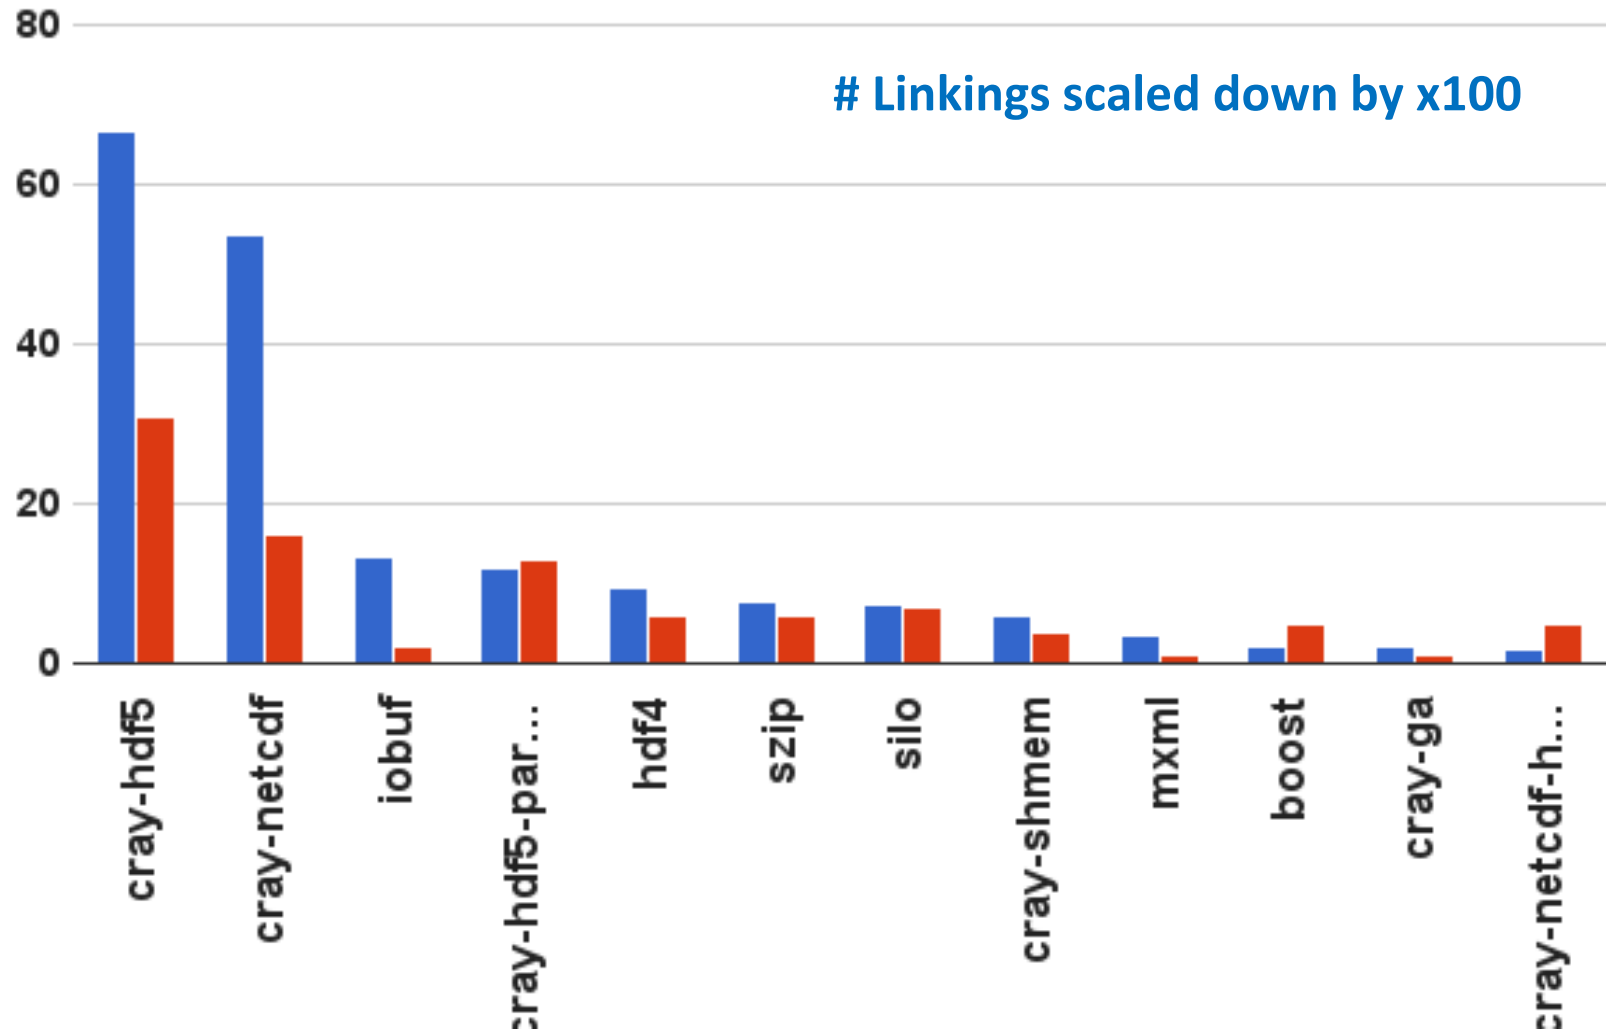
- What is “the most used” ?
 - By the number of linkings
 - By the number of unique users
- Use “module name” to identify library
 - Multiple object files may be associated with a module
 - Likely these libraries are provided via modulefile by vendor or center’s staff

Resistance to path changes as long as ReverseMap is maintained
- Script: contrib/library_usage.py

Most Used Libraries: Numerical




Most Used Libraries: Prog. & I/O



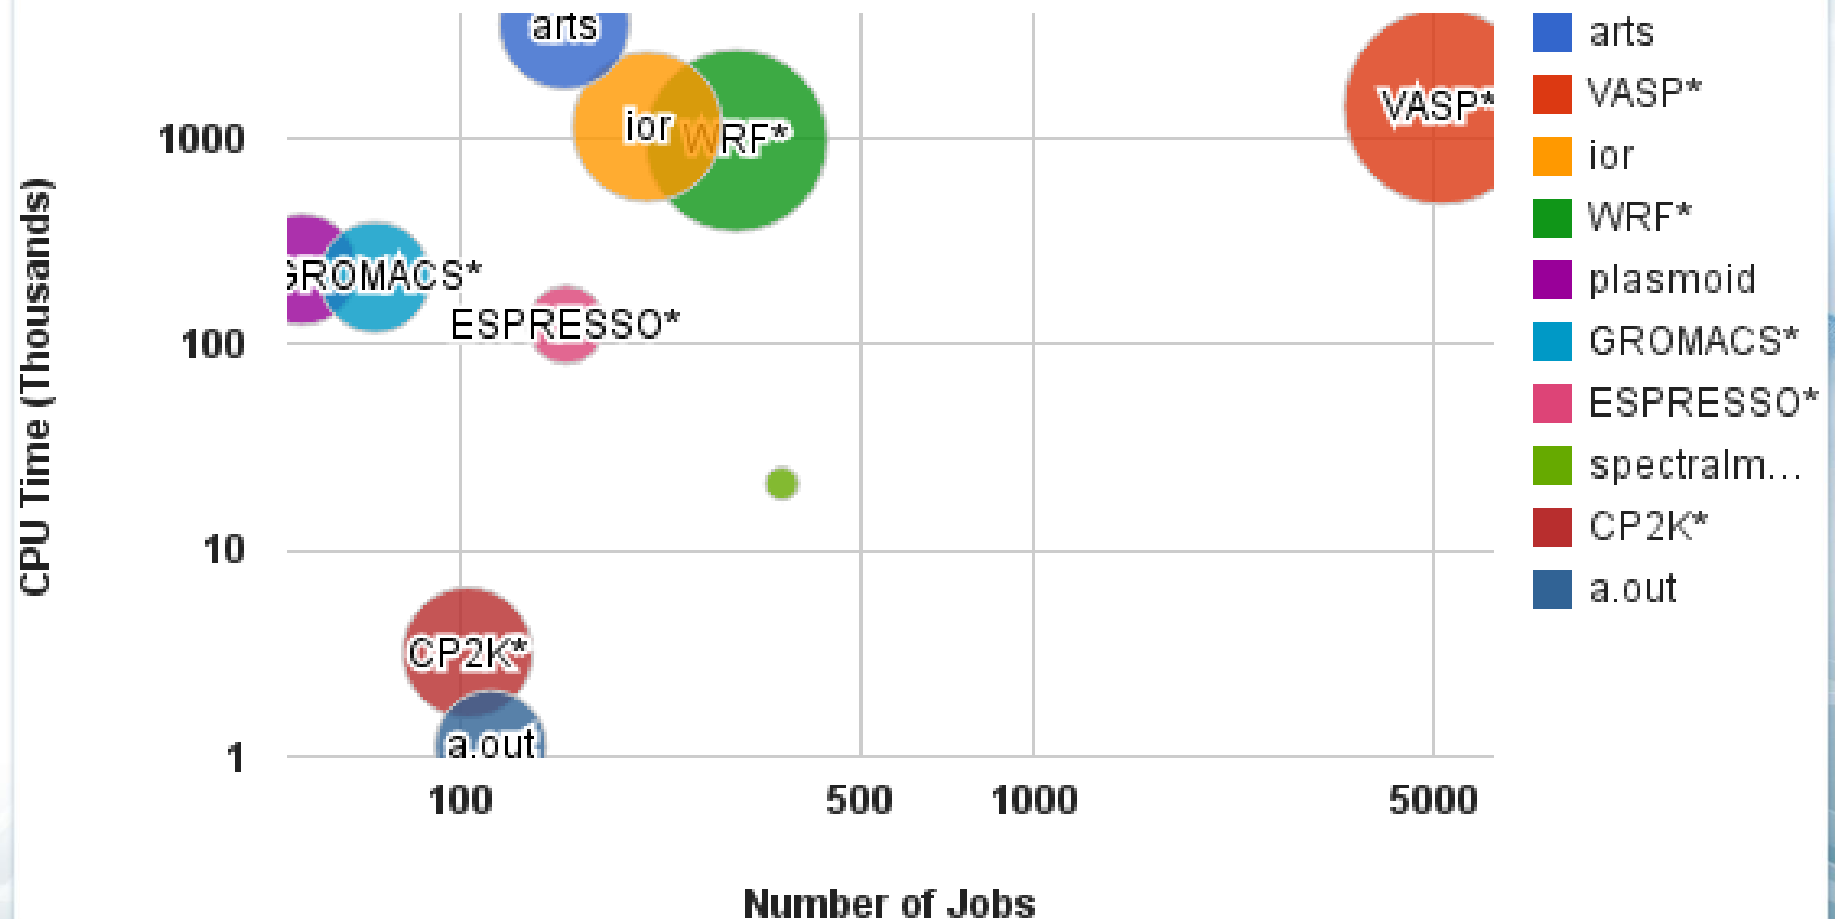
Top Executables

- Track only how much time spent by the parallel job
 - Not the entire job script
 - Can be correlated with other accounting to get the ratio of the parallel job over the entire job script
- Track the actual number of compute cores used in the parallel job
 - Done by parsing the argument given to parallel launcher
- Can show how the launched executable was built → provenance data

Top Executables



Top Executables: KAUST



Software Pruning

- How or when to remove software (version) on the system ?
 - Because newer versions are available
 - Because of lack of use
 - To free up disk space and/or support time
- XALT can provide data-driven decision
 - Show when the last time each library was used (linked against), and by whom (user)
 - Allow for targeted notification to users (to upgrade version, migrate to different library, etc)

New Functionality

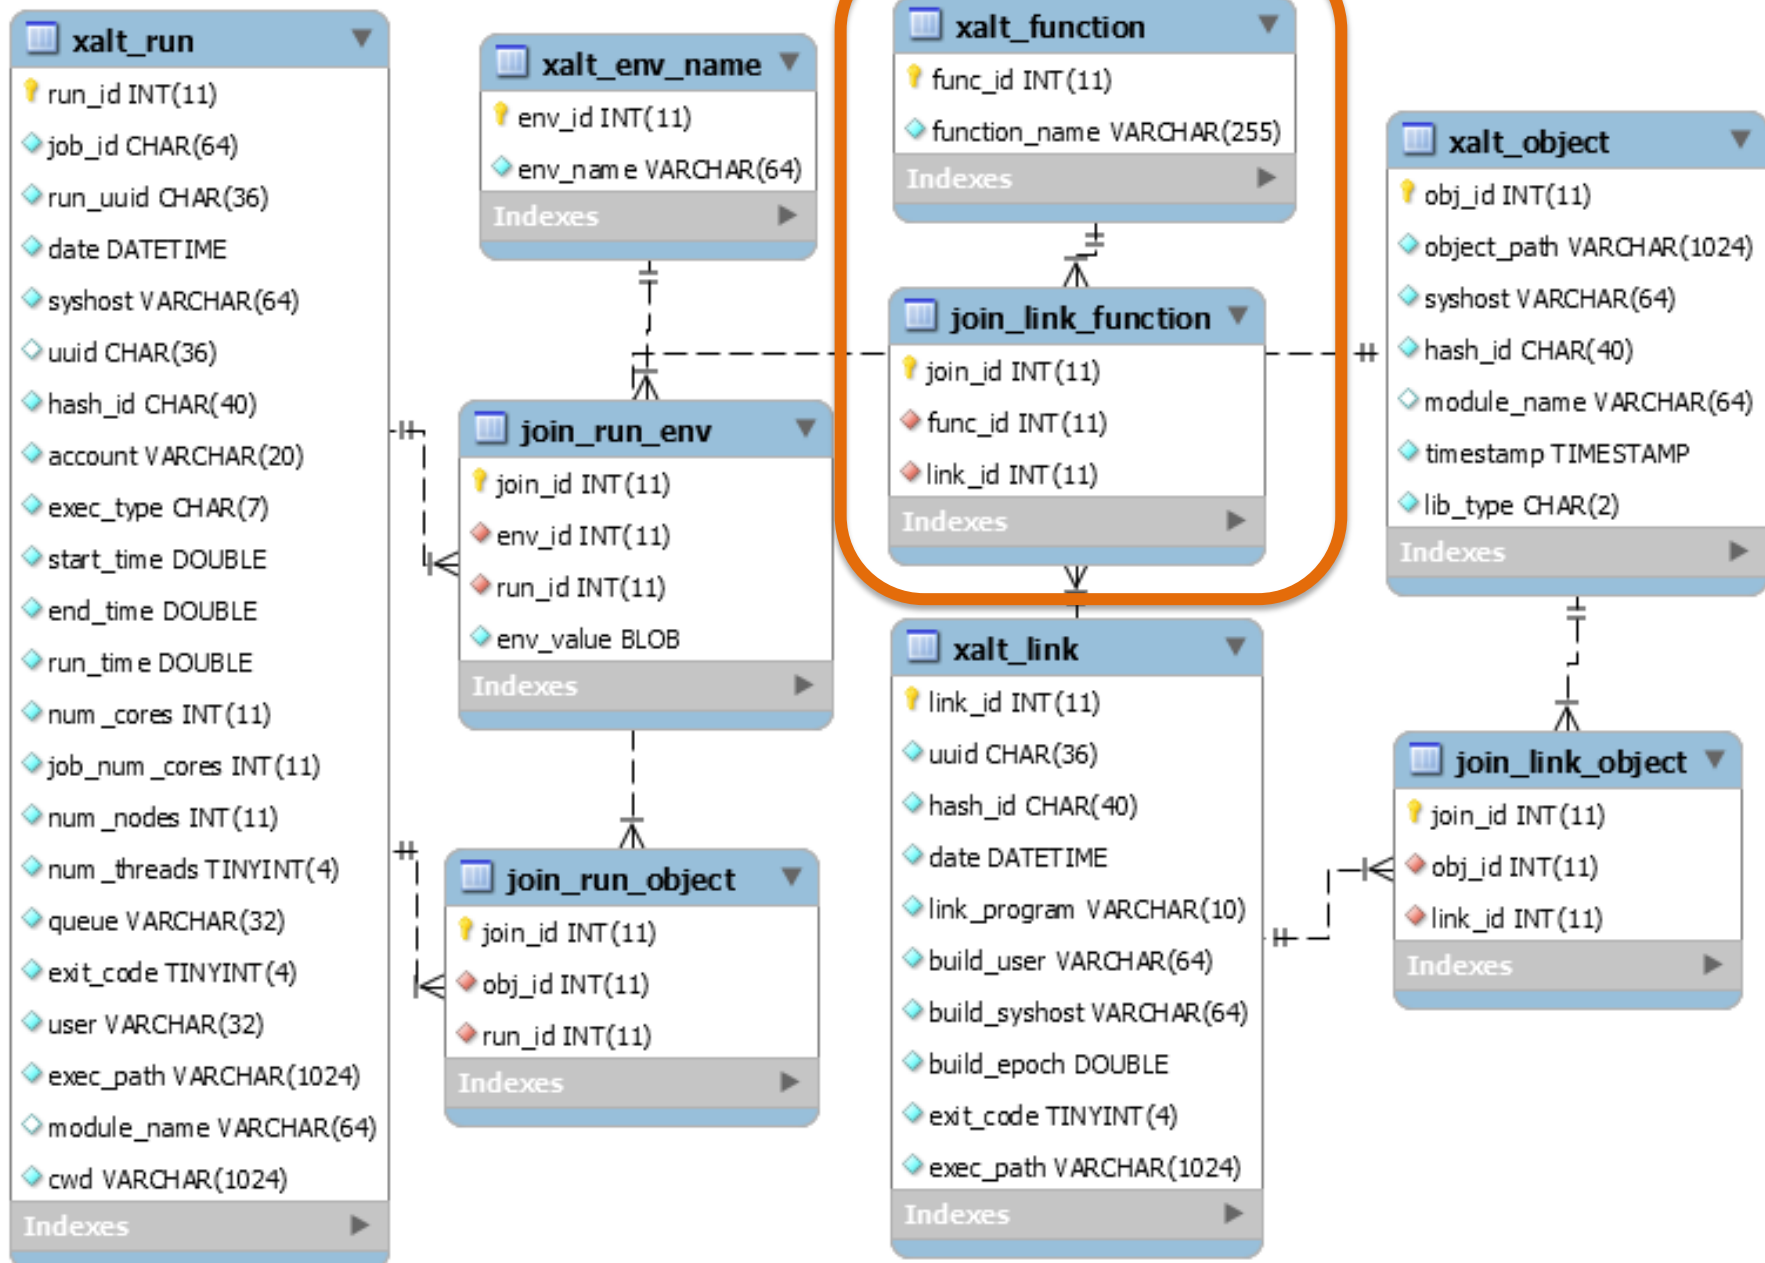
Function Tracking

- Recently added functionality (version $\geq 0.7.0$)
- Only track functions (a.k.a. subroutines / symbol names) that are resolved by external libraries
 - Does not track user defined functions
 - Does not track auxiliary functions in libraries
- Currently does not track which library resolves the functions

Although this can be done heuristically after the fact

Function Tracking (2)

- Collect the list of library / object files whose functions we are interested in tracking
 - Generated by traversing the directories of library files in modulefiles (typically used as argument to “-L” linker flag) ← already in ReverseMap file



Example Query

- Most called functions

```
SELECT trim(function_name),
       count(*)
FROM xalt_link xl,
     join_link_function lf,
     xalt_function xf
WHERE build_syshost = 'darter'
      AND xl.link_id = lf.link_id
      AND lf.func_id = xf.func_id
GROUP BY function_name
ORDER BY cnt DESC
LIMIT 100
```

FunctionName	cnt
o	4108
__intel_new_feature_proc_init	3469
std::ios_base::Init::Init()	1680
std::ios_base::Init::~Init()	1680
__gxx_personality_v0	1620
for_set_reentrancy	1450
for_rtl_finish_	1450
for_rtl_init_	1450
std::basic_ostream<char,	1414
vtable	1394
typeinfo	1238
for_write_seq_lis	1197
mpi_finalize_	1147
mpi_comm_rank_	1134
mpi_init_	1133
MPI_Comm_rank	1131
_intel_fast_memset	1120
for_write_seq_lis_xmit	1103
operator	1101
onst	1101
mpi_comm_size_	1100
_gfortran_set_args	1063
_gfortran_set_options	1063
onst&)	1063
std::cout	1052
for_open	1038
MPI_Comm_size	1030
for_write_seq_fmt	999
onst&,	996
_gfortran_st_write	993
_gfortran_st_write_done	993
std::basic_ostream<char,	970

Example Query

- BLAS' mat-mul use

```
SELECT distinct(SUBSTRING_INDEX(
exec_path,'/',-1)) as exe,
    build_user
FROM xalt_link xl,
    join_link_function lf,
    xalt_function xf
WHERE build_syshost = 'darter'
    AND xl.link_id = lf.link_id
    AND lf.func_id = xf.func_id
    AND xf.function_name
        LIKE '%gemm%'
GROUP BY exe
```

exe	build_user
Bilayer_x86_64.out	
Bilayer_x86_64_mpi.out	
Bilayer_x86_64_omp.out	
CHIMERA3D_cray	
MASS.so	
R_X11.so	
R_de.so	
arts	
average.x	
bands.x	
bgw2pw.x	
cairo.so	
class.so	
cp.x	
d3.x	
dist.x	
dmrg	
dos.x	
elk	
epsilon.x	
even-serial_a.out	
even_a.out	
fd.x	
fd_ef.x	

XALT Portal

A web interface to more easily get XALT data:

- Used by center's staff to easily get high level library, compiler, and executable usage
- From any of those “entry points”, can drill-down to users associated with library/compiler/executable, and their jobs and job environment
- Can search who uses a particular library or executable

Allow targeted notification in case of buggy library, retired versions, etc



Usage

Results presented here are taken from XALT database at National Institute for Computational Sciences at Oak Ridge National Laboratory. XALT went live on November 2014 the data presented here is realtime.rnation that most computing centers need or want.

Select syshost and date range for your queries.


Select Syshost*

darter

September 1, 2015 - October 31, 2015

Submit

Modules Usage



Modules	Count
1 alps	3777
2 boost	9
3 cce	922
4 cray-hdf5	308
5 cray-hdf5-parallel	470
6 cray-libsci	1435
7 cray-mpich	2790

Click Modules to get Version details

[Count = Number of time Object was Linked]

Further Details

List of Version(s) (for given Module)

Modules	Versions	Count
1 cray-libsci	13.0.3	1435

[Count = Number of time Object was Linked for given Module-Version]

List of User(s) (for given module-version)

Users	Earliest_LinkDate	Latest_LinkDate	Count
1	2015-09-17 17:09:18	2015-10-14 17:41:43	262
2	2015-09-17 16:44:11	2015-10-30 19:09:23	205
3	2015-09-17 00:29:32	2015-10-28 11:57:39	158
4	2015-09-01 22:33:17	2015-10-28 18:54:13	143
5	2015-09-03 13:05:52	2015-10-27 22:47:26	117
6	2015-09-01 14:26:11	2015-10-23 11:27:52	96
7	2015-09-22 14:01:31	2015-10-08 18:45:03	94

[Count = Number of Executable linked by the User]

List of Executable(s) (for given user-module-version)

Executable Name	LinkDate_Oldest	LinkDate_Latest	Count
1 a.out	2015-09-18 14:40:49	2015-10-14 17:23:34	22
2 manypw.x	2015-10-06 15:31:01	2015-10-07 13:22:46	4
3 pw.x	2015-10-06 15:31:01	2015-10-07 13:22:46	4
4 kpoints.x	2015-10-06 15:33:18	2015-10-07 13:23:56	4
5 dist.x	2015-10-06 15:33:18	2015-10-07 13:23:56	4

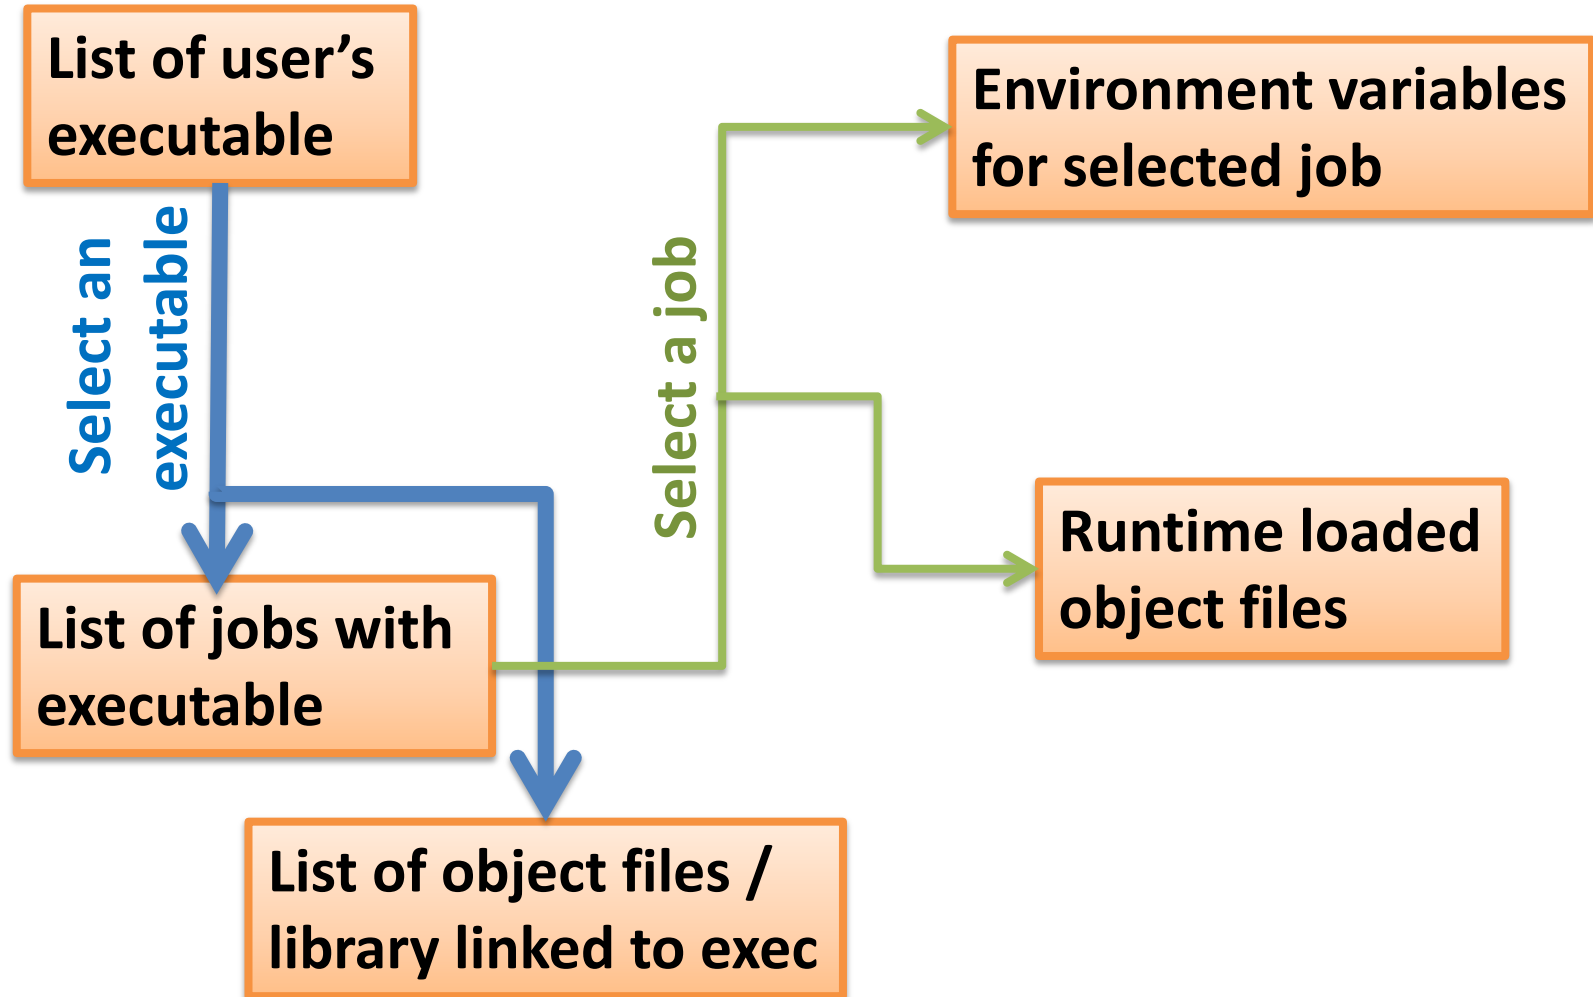
XALT Portal for User Provenance

- “How did I build my exec x months ago ?”
“What was the default MPI / compiler / libraryX at the time ?”
- Allow user to know the history and origin, i.e. “provenance”, of the software they run

Different type of users:

- Run their own executable
Run executable provided by the Center
Run executable built by another user
- Helps with reproducibility of research conducted with such software

User Provenance



User Software Provenance Get run/link details for given user

Select Syshost*

darter ▼



October 20, 2015 - November 18, 2015 ▼

Enter User ID *



Submit

Further Details

List of Executable(s)

	Executable	No_Jobs
1	ChartAlterationInteriorProper_Form_Test_Darter_Cray	12
2	ChartAlteration_IG_EP__Form_Test_Darter_GNU	12
3	ChartAlteration_IG_EP__Form_Test_Darter_Cray	11
4	ChartAlteration_IP__Form_Test_Darter_Cray	9
5	Mesh_Form_Test_Darter_Cray	8
6	ChartAlterationInteriorProper_Form_Test_Darter_GNU	8
7	ChartStream_Form_Test_Darter_GNU	6

[Count = Number of times executable was run]

Executable Path	Build Date	Link Program	Exit Code	Build User	Job Run[T/F]	Unique Id
1 /nics/d/home/aeasis_0.7/Programs/Uni tTests/Mathematics/Manifolds/Charts/Intermesh es/Executables/ChartAlterationInteriorProper_ Form_Test	2015-10-29 15:26:44	ftn_driver	0		✓	3ede0761-b90a-40fa-8c7d-ba8619f899d3
2 /nics/d/home/aeasis_0.7/Programs/Uni tTests/Mathematics/Manifolds/Charts/Intermesh es/Executables/ChartAlterationInteriorProper_ Form_Test	2015-10-29 15:23:02	ftn_driver	0		✓	68b44b30-577f-41cd-9ed3-48698107bc66
3 /nics/d/home/aeasis_0.7/Programs/Uni tTests/Mathematics/Manifolds/Charts/Intermesh es/Executables/ChartAlterationInteriorProper_ Form_Test	2015-10-29 15:20:19	ftn_driver	0		✓	c66179d6-780f-4d6b-be05-d010e9de9694
4 /nics/d/home/aeasis_0.7/Programs/Uni tTests/Mathematics/Manifolds/Charts/Intermesh es/Executables/ChartAlterationInteriorProper_ Form_Test	2015-10-29 15:17:18	ftn_driver	0		✓	ac247ba6-ad26-492a-8db2-3b18fbf6e44b
5 /nics/d/home/aeasis_0.7/Programs/Uni tTests/Mathematics/Manifolds/Charts/Intermesh es/Executables/ChartAlterationInteriorProper_ Form_Test	2015-10-29 15:14:59	ftn_driver	0		✓	ca2dcba4-f456-4a69-aa7d-29ef85301806
6 /nics/d/home/aeasis_0.7/Programs/Uni tTests/Mathematics/Manifolds/Charts/Intermesh es/Executables/ChartAlterationInteriorProper_ Form_Test	2015-10-29 15:01:57	ftn_driver	0		✓	eff23a99-94d4-44b0-9e17-cabefdddeb30d
/nics/d/home/cardall/aeasis_0.7/Programs/Uni tTests/Mathematics/Manifolds/Charts/Intermesh es/Executables/ChartAlterationInteriorProper	2015-10-29					76b97080-ecf4-45f6-81ed-

Objects Linked (to the given Executable)

Object Path	Module Name	Object Date	Object Type
1 /opt/cray/xpmmem/0.1-2.0502.55507.3.2.ari/lib64/libxpmmem.a	xpmmem/0.1-2.0502.55507.3.2.ari	2015-04-10 15:54:05	a
2 /opt/cray/wlm_detect/1.0-1.0502.53341.1.1.ari/lib64/libwlm_detect.a	wlm_detect/1.0-1.0502.53341.1.1.ari	2015-04-10 15:54:05	a
3 /opt/cray/ugni/5.0-1.0502.9685.4.24.ari/lib64/libugni.a	ugni/5.0-1.0502.9685.4.24.ari	2015-04-10 15:54:05	a
4 /opt/cray/udreg/2.3.2-1.0502.9275.1.12.ari/lib64/libudreg.a	udreg/2.3.2-1.0502.9275.1.12.ari	2015-04-10 15:54:05	a
5 /nics/e/sw/xlc30_cle5.2_pe2014-09/silo/4.9.1/cle5.2_gnu4.9.1/lib/libsiloh5.a	silo/4.9.1	2015-04-10 15:54:04	a
6 /opt/cray/rca/1.0.0-2.0502.53711.3.127.ari/lib64/librca.a	rca/1.0.0-2.0502.53711.3.127.ari	2015-04-10 15:54:05	a
/opt/cray/dmi/5.0.6-1.0000.10439.140.2.ari/lib			

Job Run Details (nC-nJC-nN-nT ~ #Cores-#JobMunkCores-#Nodes-#Threads)

RunId	JobId	Run Date	nC-nJC-nN-nT	Account	Exec Type	Run Time (sec)	ExitCode	Run User	CurrentWorkingDir
1 535974	660918 [REDACTED] edu	2015-10-29 15:24:34	1 64 1 0	[REDACTED]	binary	5.56	0	[REDACTED]	/lustre/medusa/[REDACTED]7/Programs/UnitTests/Mathematics/Manifolds/Charts/Intermeshes/Executables

Run Environment Details (for the given Job)

Environment Variable	Value
1 ALT_LINKER	/sw/x86_64/gcc/4.8.2/bin/ld
2 ASSEMBLER_X86_64	/opt/cray/cce/8.3.9/cray-binutils/x86_64-unknown-linux-gnu/bin/as
3 ATP_HOME	/opt/cray/atp/1.8.0
4 ATP_MRNET_COMM_PATH	/opt/cray/atp/1.8.0/libexec/atp_mnet_commnode_wrapper
5 ATP_POST_LINK_OPTS	-Wl,-L/opt/cray/atp/1.8.0/libApp/
6 CC_X86_64	/opt/cray/cce/8.3.9/CC/x86-64
7 CPU	x86_64
8 CRAYLIBS_X86_64	/opt/cray/cce/8.3.9/craylibs/x86-64
9 CRAYLMD_LICENSE_FILE	/opt/cray/cce/cce.lic
10 CRAYOS_VERSION	5.2.40
11 CRAYPE_DIR	/opt/cray/craype/2.2.1
12 CRAYPE_NETWORK_TARGET	sdsl

Conclusions

- XALT has been in production for over a year
- XALT has been successfully deployed on multiple HPC centers to support their operations
- XALT helps stakeholders make data-driven decision on software support
- Further analysis on XALT data may yield more understanding of interesting users' behavior
- Source: **<https://github.com/Fahey-McLay/xalt>**

Acknowledgment

- This work was supported by the NSF award 1339690 entitled “Collaborative Research: SI2-SSE: XALT: Understanding the Software Needs of High End Computer Users.”
- Thanks to the XALT community for feedback and bug reports