



Automatically Encapsulating HPC Best Practices Into Data Transfers

Paul Z. Kolano

NASA Advanced Supercomputing Division

paul.kolano@nasa.gov

Outline of Presentation



- **Introduction**
- Transport tuning and selection
- Global resource management
- File system optimization
- Conclusions

Introduction



- Data transfers are part of life in HPC environments
 - Finite storage capacity
 - Transfer to cheaper tape storage
 - Back up existing data
 - Make room for new data
 - Transfer from tape storage to reprocess old data
 - Transfer between file systems to fix imbalances
 - Finite computational capacity
 - Transfer from off-site systems with cheaper pre-processing
 - Transfer to off-site systems for cheaper post-processing

Introduction (cont.)



- User transfer concerns
 - Ease to use, integrity, turnaround time
- Administrator and owner transfer concerns
 - Environment stability, cost effectiveness
- These items can conflict with each other
 - Easy to use tools or those ensuring integrity may not be fast
 - Easiest file structure may degrade tape performance
 - Fastest turnaround time may lead to resource exhaustion
- Takes HPC expert to reconcile conflicts
 - Understands and applies accepted best practices to achieve fast and efficient verified transfers without impact on stability

Goal




- Let scientists focus on science without wading through documentation on transfer best practices
 - Specify transfers in simplest, naive fashion
 - Source and destination
- Provide tool to perform transfer as if scientist were HPC expert
 - Choose appropriate tools and optimize for best performance
 - Fully utilize available resources without starving other users
 - Manage files appropriately by file system type to ensure efficient access by later and/or behind the scenes processes

Shift: Self-Healing Independent File Transfer

- Satisfies user requirements
 - Simple cp/scp syntax for local/remote transfers
 - End-to-end integrity via checksums and sanity checks
 - High speed via transport selection/tuning and automatic parallelization
- Satisfies administrator and owner requirements
 - Helps prevent resource exhaustion that leads to environment instability
 - Global throttling to allocate resources fairly
 - Load balancing to avoid highly loaded hosts
 - Automatic striping to avoid imbalanced disk utilization
 - Helps prevent wasted resources that impact cost effectiveness
 - Allows easy utilization of idle resources
 - Reduces wasted CPU cycles during jobs due to inefficient disk I/O
 - Prevents issues leading to inefficient tape I/O

Shift: Self-Healing Independent File Transfer (cont.)



- Used in production at NASA's Advanced Supercomputing Facility for over 3.5 years
 - User transfers across local/LAN/WAN
 - Disaster recovery backups to/from remote organizations
 - Rebalancing entire multi-PB Lustre file systems
 - etc.
- Facilitated transfers of over 14 PB in past year
 - 8 PB local transfers
 - 4 PB LAN transfers
 - 2 PB WAN transfers
- "I used to hate archiving data - now I almost look for a reason to archive something" –Shift user

Shift Interface



```
> shiftc --create-tar /nobackup/user1/dataset1 archive1:dataset1.tar
```

Shift id is 36

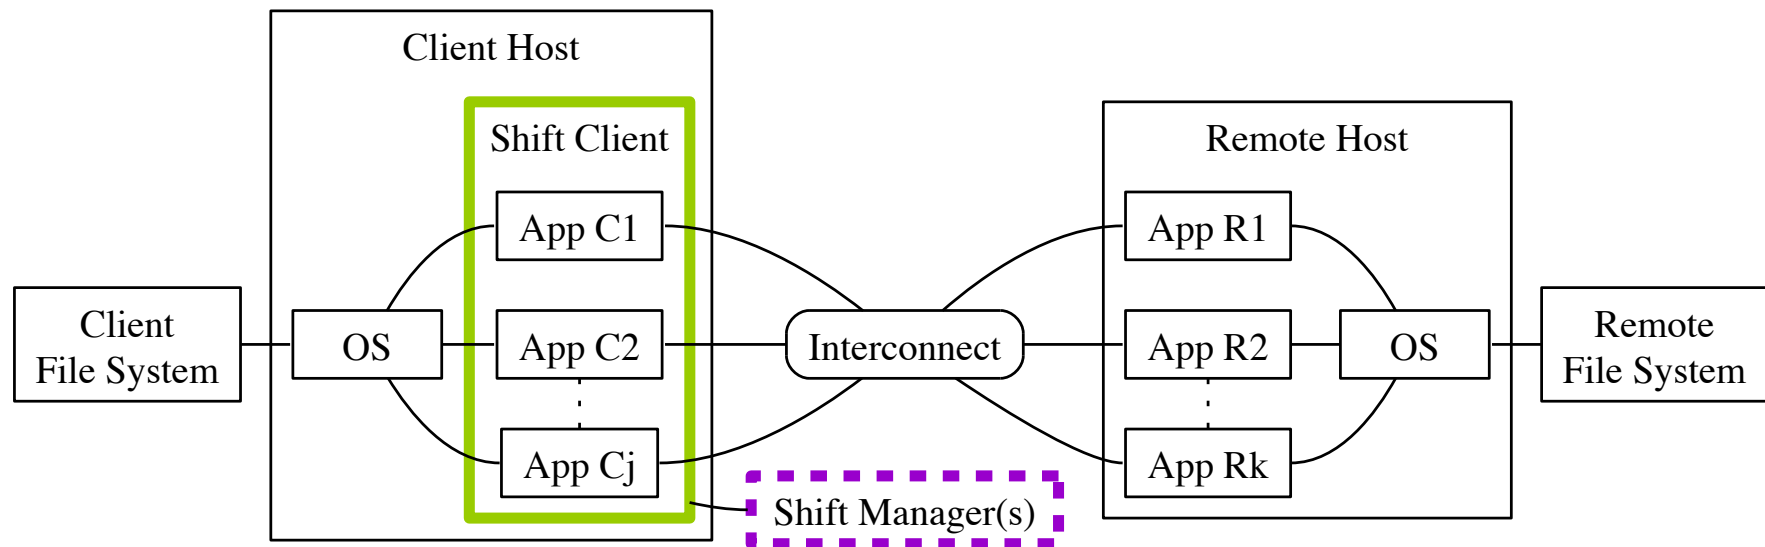
Detaching process (use --status option to monitor progress)

```
> shiftc --status
```

id	state	dirs	files	file size	date	run	rate
		sums	attrs	sum size	time	left	
-----+-----+-----+-----+-----+-----+-----+-----							
34	error	0/0	23121/23121	39.5TB/39.5TB	10/02	2d14h32m5s	175MB/s
		46222/46242	23111/23121	79TB/79TB	10:26		
35	done	1/1	5131/5131	303GB/303GB	10/05	1m35s	3.19GB/s
		10262/10262	5132/5132	605GB/605GB	12:28		
36	run	24/24	26656/26656	1.78TB/1.78TB	10/06	2h48m37s	176MB/s
		15463/53312	10/26684	1.02TB/3.56TB	12:11	1h47m55s	

Shift Components

- Command-line client
 - Performs file operations and reports results to manager
- Command-line manager
 - Invoked by clients to track operations and parcel out work



Outline of Presentation



- Introduction
- **Transport tuning and selection**
- Global resource management
- File system optimization
- Conclusions

Transport Tuning and Selection



- Shift includes built-in local/remote transports and checksum capabilities
 - Fully functional out of the box
 - Perl-based equivalents of cp, sftp, fish, m(d5)sum
- Shift calls higher performance tools when available
 - bbcp, bbftp, gridftp, mcp, rsync, msum
 - Knows how to construct command-lines and parse output
- Tune transports for optimal performance
- Select transports based on transfer characteristics

Transport Tuning



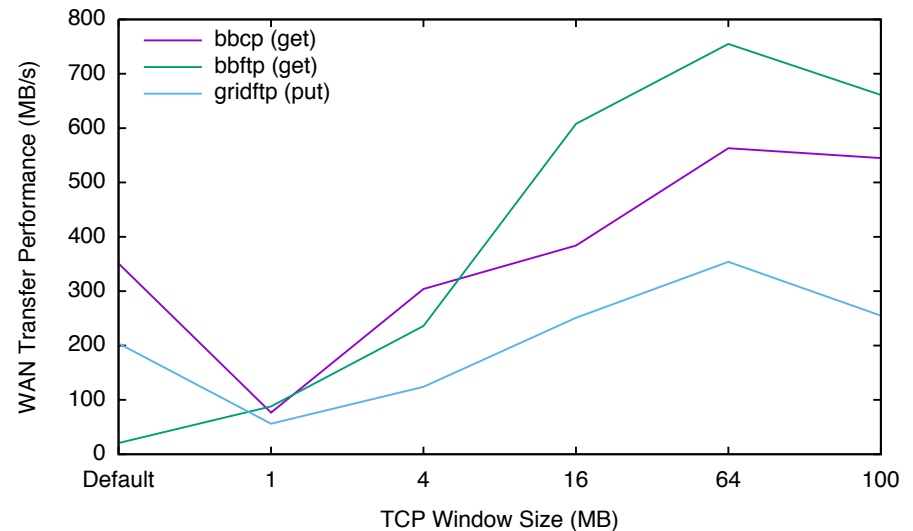
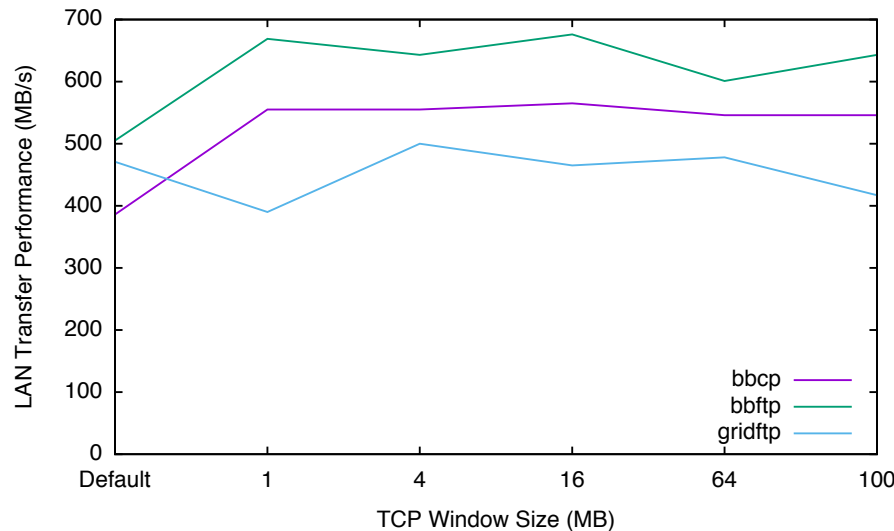
- TCP-based transports
 - bbcp, bbftp, gridftp
 - Choose TCP window size
- Transports with internal parallelism
 - TCP streams (bbcp, bbftp, gridftp) or threads (mcp, msum)
 - Choose appropriate level of parallelism
- SSH-based transports
 - fish, rsync, sftp-perl
 - Choose fastest SSH cipher and MAC algorithm

TCP Window Size Tuning



- TCP window is amount of data sender or receiver willing to buffer while waiting for acknowledgment
- Optimal value is bandwidth delay product (BDP)
 - $\text{bandwidth} * \text{round-trip time}$
- Constrained by configured operating system limits
 - e.g. Linux `net.core.[wr]mem_max`
 - Single stream only achieves bandwidth if limit at least BDP

TCP Window Size Tuning (cont.)



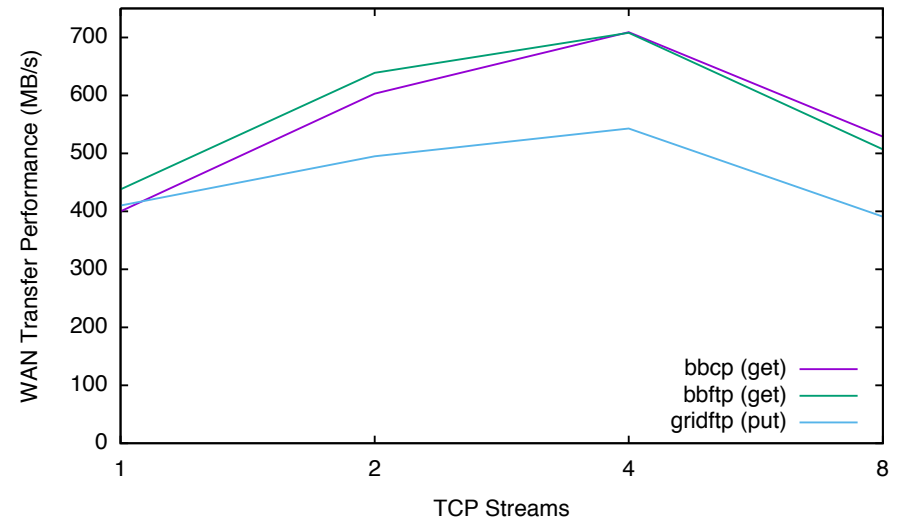

- Shift determines latency using icmp/echo/syn ping
- Shift guesses bandwidth based on network type and client hardware if not given via --bandwidth
 - Bandwidth difficult to compute a priori
- Chooses window size up to operating system limit

Transport Parallelism Tuning



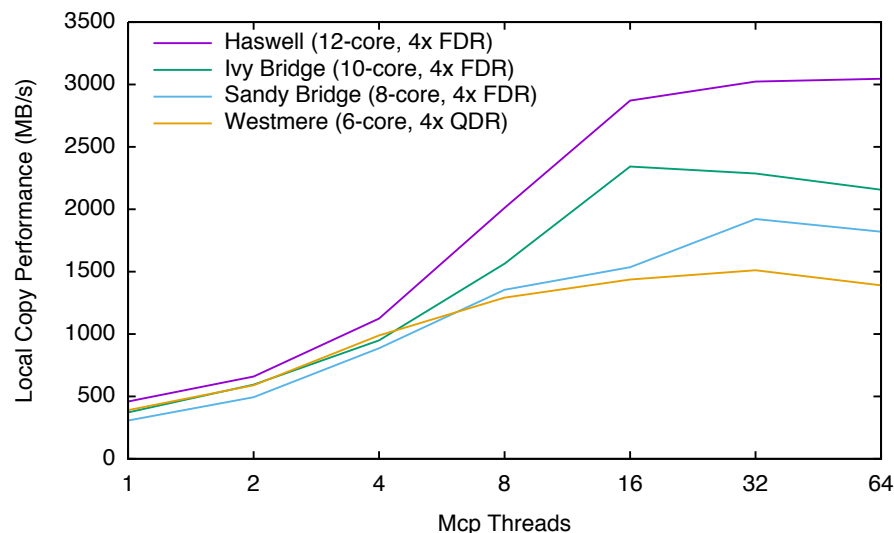

- Number of streams in TCP-based transports
 - Overcome improperly configured TCP window maximums
 - Overcome improperly specified TCP window
 - Overcome interference by cross traffic
- Number of threads in mcp and msum
 - Take advantage of excess resource capacity on one host

Transport Parallelism Tuning (cont.)



- Shift chooses streams based on bandwidth available beyond operating system window limit
- A minimum value can be configured for LAN/WAN to help overcome cross traffic

Transport Parallelism Tuning (cont.)



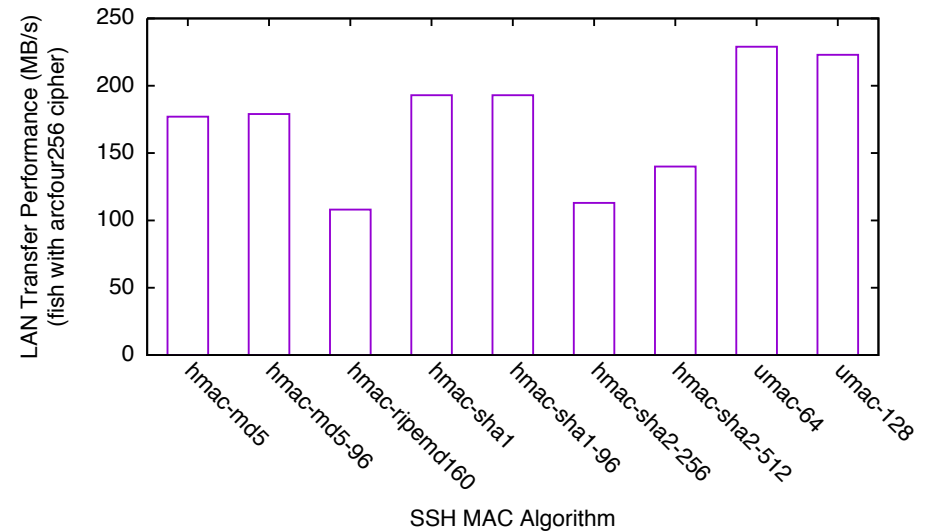
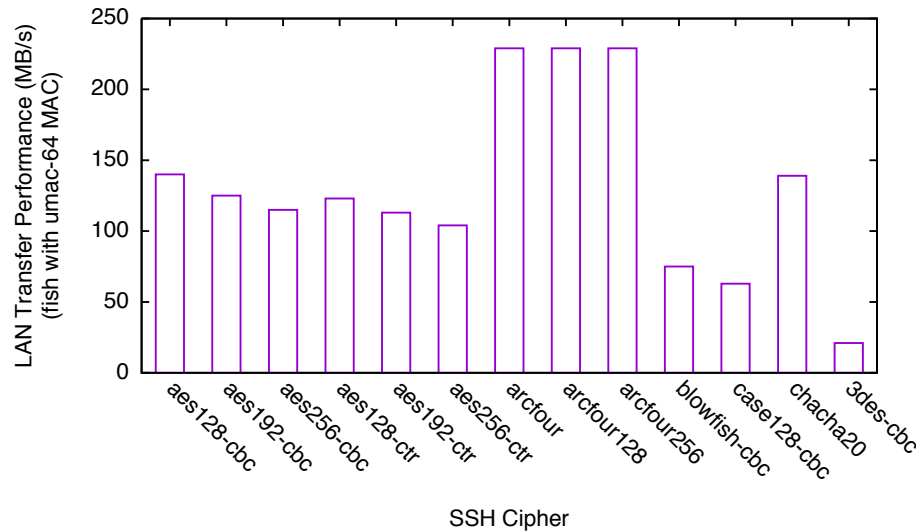
- Threads can be centrally configured on the manager
- High thread counts can induce high load on shared resources
 - Intentionally set lower than optimal at NAS due to high load on archive front-ends

SSH Cipher and MAC Algorithm Tuning



- SSH-based transports use SSH pipe to communicate
 - Performance directly correlated to SSH performance
- SSH does not expose TCP window settings
 - HPN SSH patches can be used for better window handling
- Main SSH tuning parameters available
 - Encryption algorithm
 - Message authentication code (MAC) algorithm

SSH Cipher and MAC Algorithm Tuning (cont.)




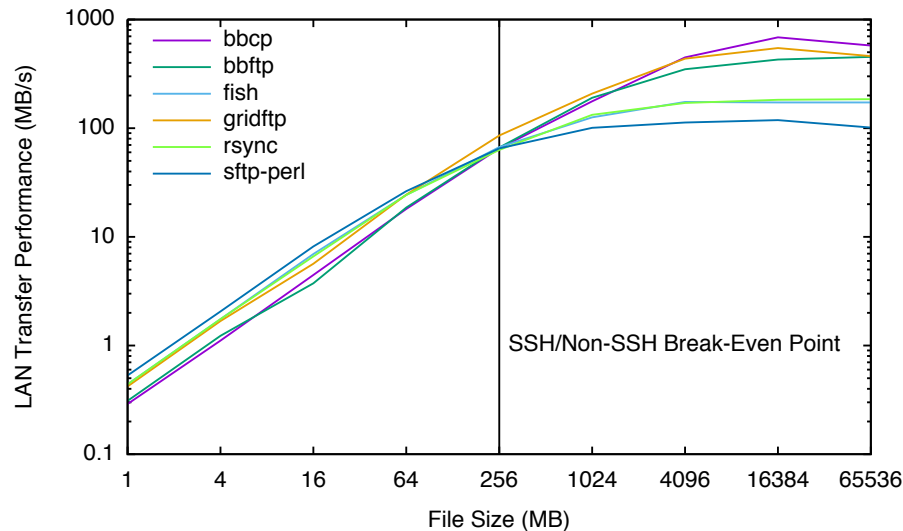
- Shift allows preferred cipher/mac order to be centrally configured
- Availability checked on client host before transfer

Transport Selection



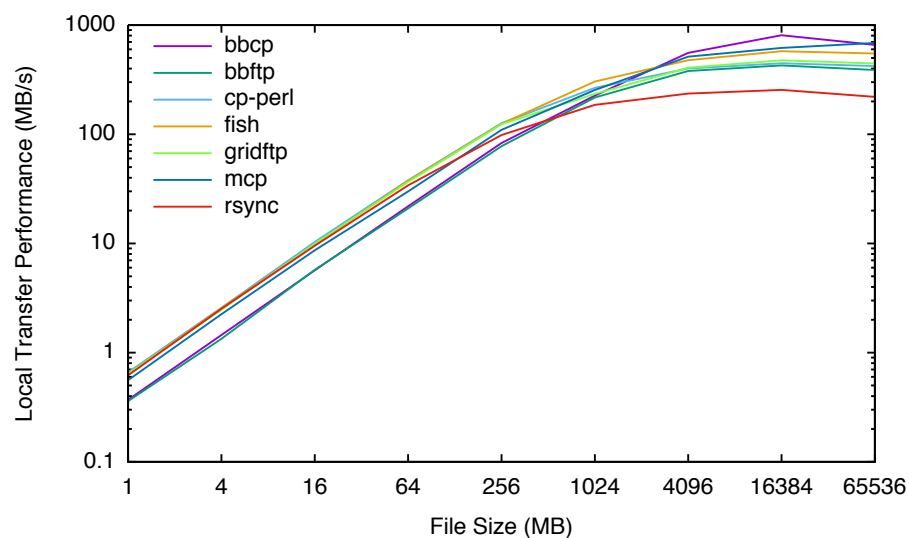
- Different transports have different strengths and weaknesses
- Supporting multiple transports provides opportunity to vary transport across each batch of files within transfer

Transport Selection (cont.)



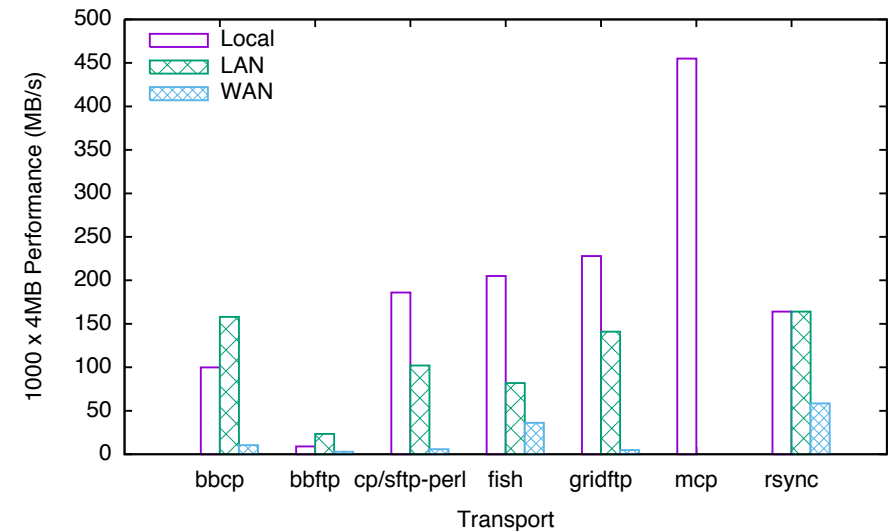

- Using a single transport does not achieve maximum performance
- Shift's support of multiple transports allows it to use the optimum transport for each batch of files

Transport Selection (cont.)



- Traditionally remote transports also perform well for local transfers

Transport Selection (cont.)



- Shift has configurable small file thresholds
 - Preferred local/LAN/WAN transports above/below
- Transport chosen using avg. size of each batch

Outline of Presentation



- Introduction
- Transport tuning and selection
- **Global resource management**
- File system optimization
- Conclusions

Global Resource Management



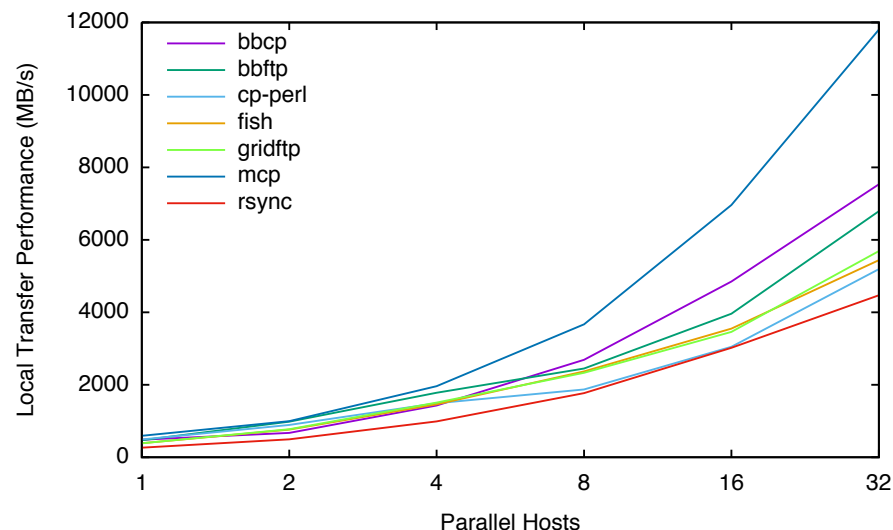
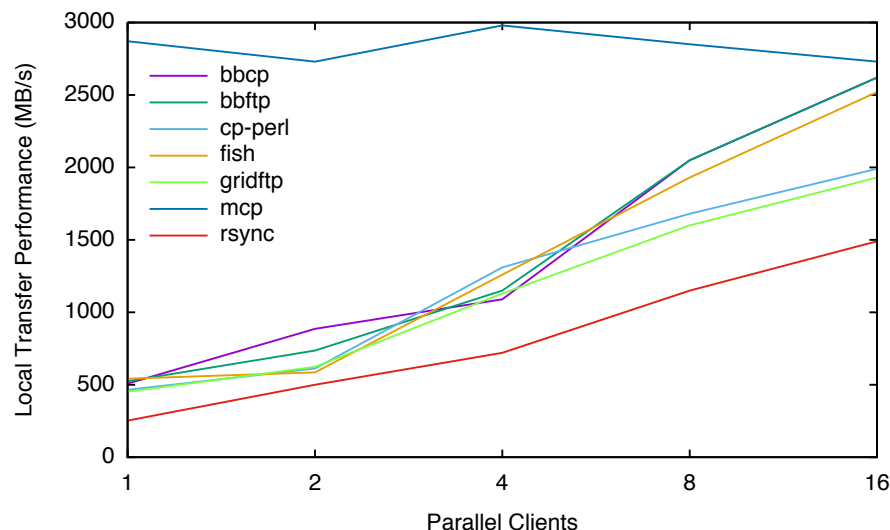
- Transfer parallelization
 - The only way to maximize performance in HPC environments is to use multiple resources at once
- Global throttling
 - If everyone tries to run at the maximum rate possible, everyone loses
- Load balancing
 - Avoid resource exhaustion on individual hosts

Transfer Parallelization



- Single host may have excess capacity
 - Transport does not have its own parallelism
 - Single client cannot fully utilize host's resources
- Full HPC environment may have excess capacity
 - Single system bottlenecks
 - Aggregate resources of many hosts
- Two complementary forms of parallelization
 - Multiple clients running on a single host
 - One or more clients running on multiple hosts

Transfer Parallelization (cont.)

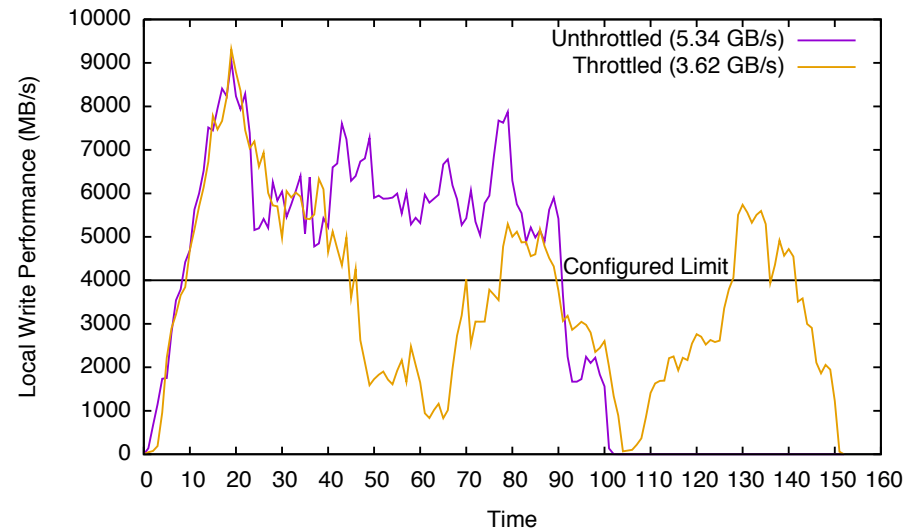


- Shift derives file system equivalency from mount information to determine parallelization candidates
- Shift can automatically parallelize all transfers using centrally configured clients/hosts values
- Can trivially run across allocated compute nodes
 - e.g. `--host-file=$PBS_NODEFILE`

Global Throttling



- Support transfers at full speed when excess capacity exists
- Prevent resource exhaustion when systems are busy

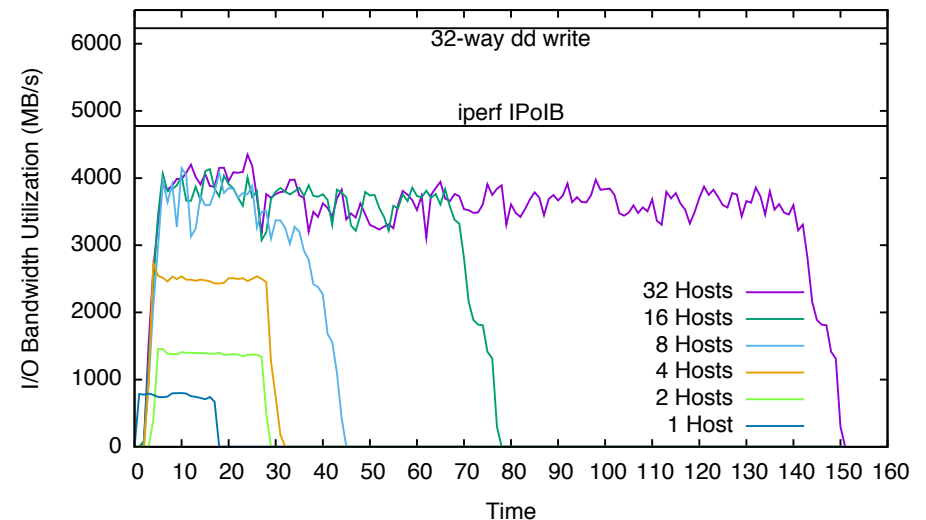


- Shift supports throttling limits on CPU %, disk usage, I/O rate, and network rate
 - Limits can be global, per user, per host, per file system
- Transfers directed to sleep until average reaches (transfer's fair share of) user's fair share of limit


Load Balancing



- Single host resource limitations impact maximum transfer speed of all processes on that host
- Less loaded hosts have more resources available for new transfers



Load Balancing (cont.)




- Shift load balances during host parallelization
 - Picks least loaded hosts when spawning clients
- Shift load balances during remote transfers
 - Remote host switched to least loaded
 - Overlap prevented between parallel clients

Outline of Presentation



- Introduction
- Transport tuning and selection
- Global resource management
- **File system optimization**
- Conclusions

File System Optimization



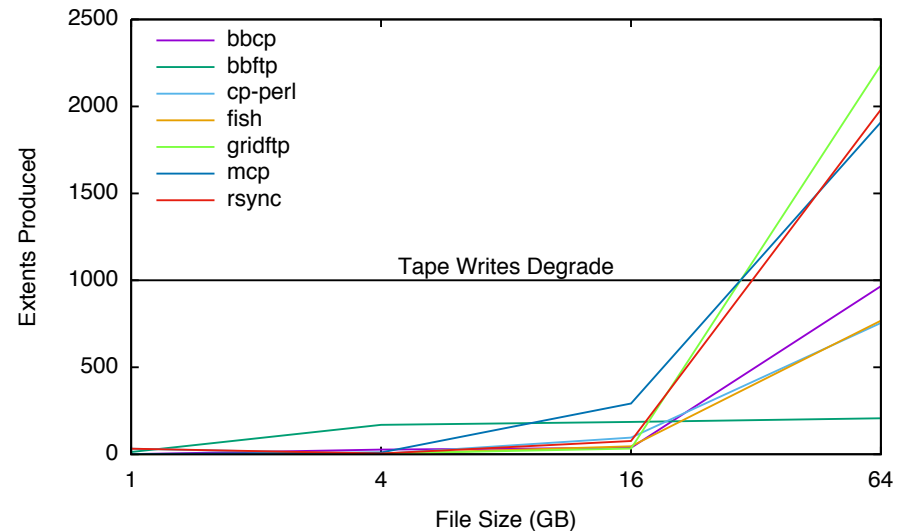
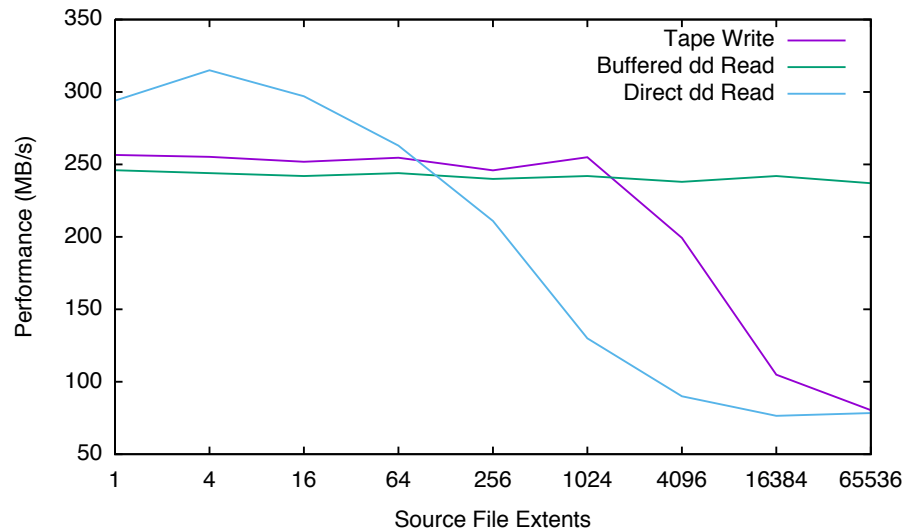
- Different file systems have different idiosyncrasies
 - Can impact performance, stability, and operating cost
 - Not always directly visible to users
- Tape optimization
 - File extents impact tape write speed
 - Sequential retrieval results in inefficient tape movement
- Tar creation/extraction
 - Tape I/O is inefficient with small files
- Lustre striping
 - Too few stripes creates later I/O inefficiencies

Tape Optimization



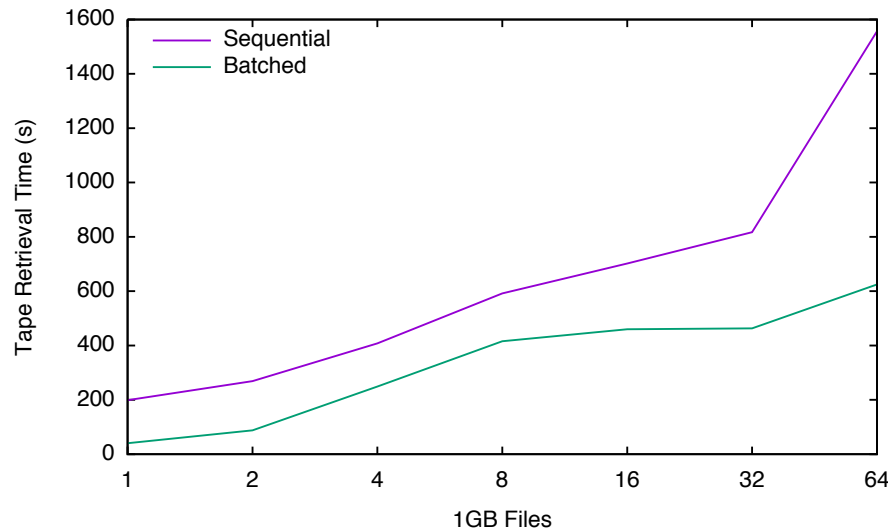
- Tape-backed file systems have limited parallelism due to large slow-moving physical components
- Inefficient access by one user can have major impact on all others
 - Large number of file extents degrades write speed
 - Sequential file retrieval inhibits minimization of internal tape movement

Tape Optimization (cont.)



- Shift can be configured to preallocate files below a given sparsity
 - Minimize extents for most common regular files
 - Minimize disk usage for large sparse files

Tape Optimization (cont.)



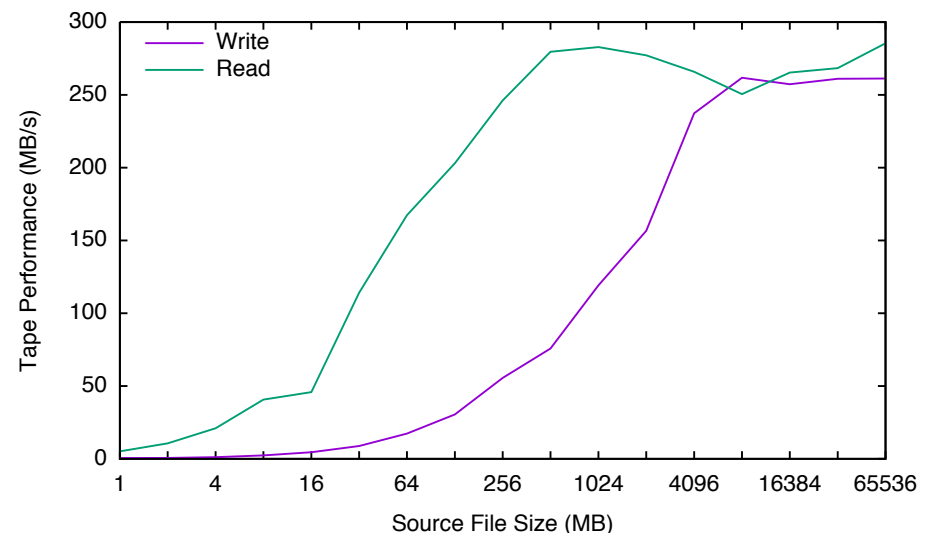
- Files automatically retrieved from tape if offline when accessed
- May be accessed in different order than stored on tape

- Shift automatically initiates a retrieval of all source files on SGI DMF file systems
- Retrieval initiated again for files in each batch in case files pushed offline before being transferred

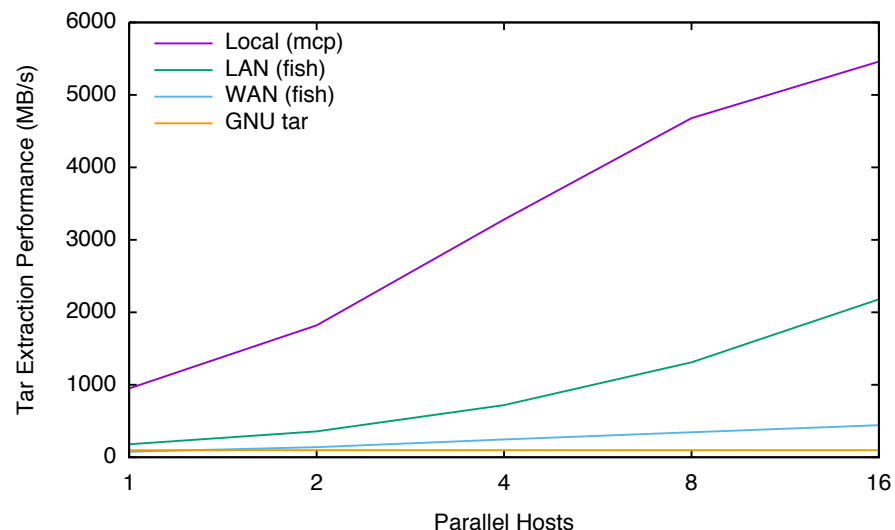
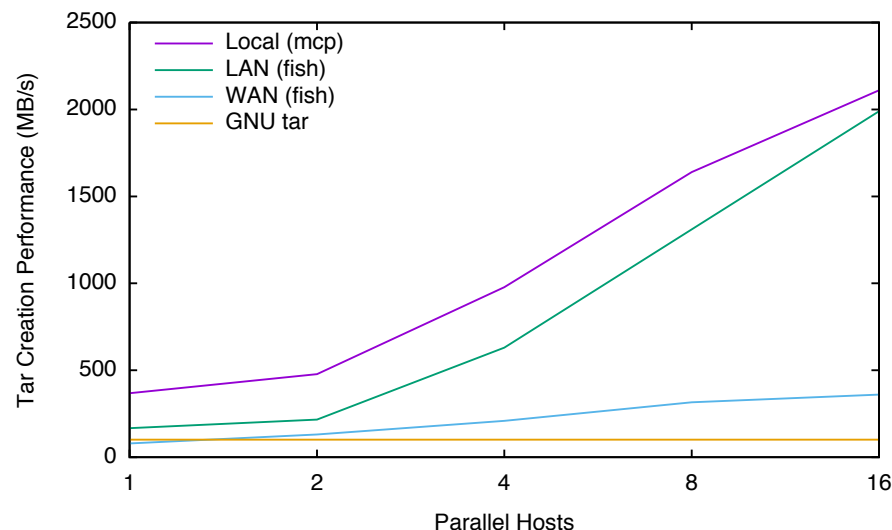
Tar Creation/Extraction



- Users prefer archiving data in original hierarchy
 - Each file accessed as needed
- Small files impact stability
 - Not migrated below certain size so consume limited disk
 - Tape I/O is inefficient
- Normally use tar files
 - Tar is very slow (100 MB/s)
 - Must retrieve to view contents
 - No assurance of integrity
 - Retrieve entire tar for one file
 - May not have quota to create



Tar Creation/Extraction (cont.)




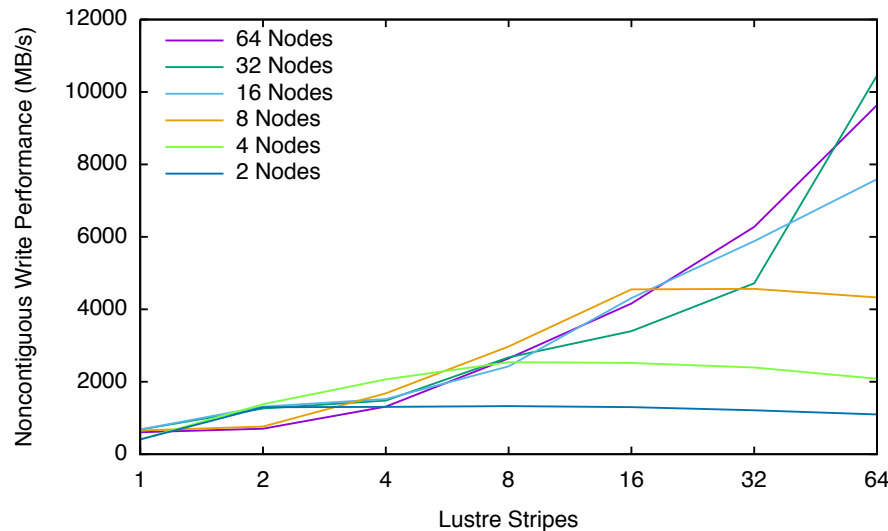

- Shift has built-in tar creation/extraction
 - Uses high performance transports and parallelism
 - Automatic creation of index files to see contents
 - Integrity verification of every file added/extracted
 - Automatic split of tar files at configurable size
 - Direct creation/extraction over network without use of quota

Lustre Striping



- Striping must be set before a file is first written
- Higher stripe count
 - More I/O bandwidth available for large files
 - More balanced distribution of large files over OSTs
- Lower stripe count
 - Less contention during metadata operations
 - Less wasted space for small files
- Striping can only be changed by copying file

Lustre Striping (cont.)



- Shift automatically stripes files according to size
 - Increases write performance during parallel transfers
 - Increases read performance during later job access
 - Reduces wasted CPU cycles due to I/O
- Preserves non-default striping when applicable

Conclusion



- Shift is an automated transfer tool that encapsulates HPC best practices to achieve better performance while preserving stability of HPC environments
 - Centralized configuration in manager component allows policies to be changed globally across all clients
 - New best practices can be incorporated transparently without user in the loop
- Shift is open source and available for download
 - <http://shiftc.sourceforge.net>
- Contact info
 - paul.kolano@nasa.gov
- Questions?